

---

---

# **MLDS 2017 Spring**

## **HW3 - Generative Adversarial Networks**

— [mldsntu2017@gmail.com](mailto:mldsntu2017@gmail.com) —

---


---

# Outline


- Task Introduction
  - Text2image generation
  - Dataset collection
- Model
  - Conditional GAN
  - Tips for training
    - Discriminator loss function
    - Objective function
- Submission and grading

# Task Introduction - text2image generation


**an all black bird**




**Bird  
Generative Model**



**this flower is  
white and pink**








**Flower  
Generative Model**



# Task Introduction - text2image generation

**a man with no hair  
and one red strong fist**

**a girl with blue hair,  
blue eyes and  
twin ponytail**




# Data Collection

## Tags

? aikatsu! 98  
? hikami sumire 9  
? hiten goane ryu 64  
? clouds 7859  
? flowers 9334  
? leaves 1778  
? long hair 54717  
? purple eyes 10468  
? purple hair 8210  
? rainbow 302  
? sky 9537  
? twintails 18309  
? umbrella 2239  
? water 9675  
? wet 5730

Not all tags are useful

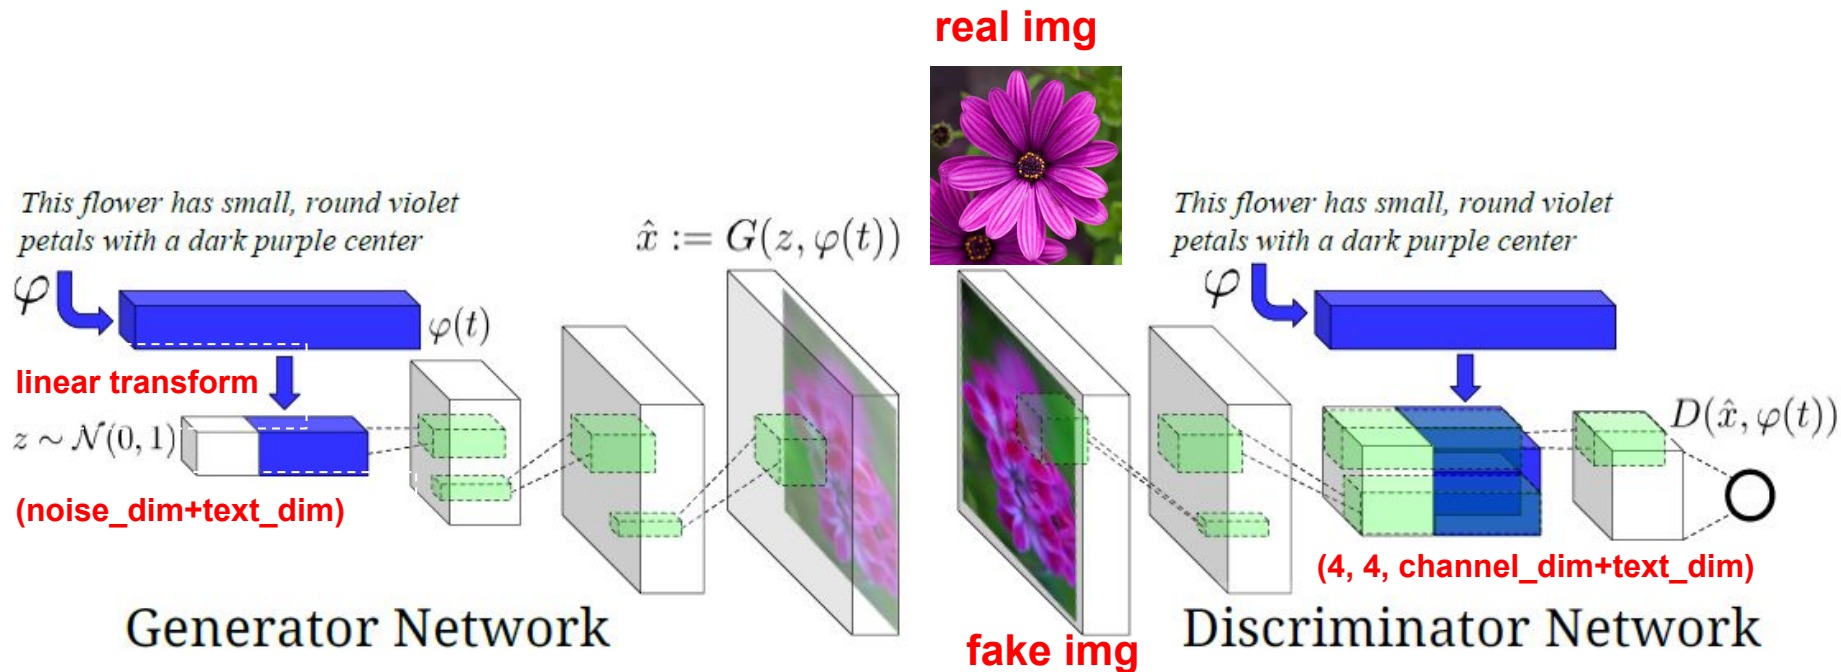


[http://konachan.net/post/show/239400/aikatsu-clouds-flowers-hikami\\_sumire-hiten\\_goane\\_r](http://konachan.net/post/show/239400/aikatsu-clouds-flowers-hikami_sumire-hiten_goane_r)

感謝樊恩宇助教蒐集data

# Model and training tips

# Conditional GAN for text2image generation



# Details for training

- Updates between Generator and Discriminator
  - 1 : 1 or 2 : 1
- ADAM with  $lr = 0.0002$ , momentum = 0.5
- gaussian or uniform noise dim = 100
- batch size = 64
- epoch = 600
  - Practically, epoch = 300 is enough



# Text feature process tool - Skip-thought vector

```
In [2]: import skipthoughts

In [3]: model = skipthoughts.load_model()
Loading model parameters...
Compiling encoders...
Loading tables...
Packing up...

In [4]: vecs = skipthoughts.encode(model, ['i love machine learning', 'text to image generation yeah'])
4
5

In [5]: vecs.shape
Out[5]: (2, 4800)
```

No matter which tool you use to process text input, please make sure you include that pre-trained model in your repository to let us run your code successfully.

skip-thought source code: <https://github.com/ryankiros/skip-thoughts>

# Image process tool - skimage and scipy.misc

```
In [35]: # convert img to tensor
In [36]: import skimage
In [37]: import skimage.io
In [38]: img = skimage.io.imread('sample.jpg')
In [39]: # resize img
In [40]: import skimage.transform
In [41]: img_resized = skimage.transform.resize(img, (64, 64))
In [42]: img.shape
Out[42]: (96, 96, 3)
In [43]: img_resized.shape
Out[43]: (64, 64, 3)
```

Install:

- `sudo apt-get install python-skimage`
- `sudo pip install --user numpy scipy`

```
In [48]: # convert tensor to img
In [49]: import scipy.misc
In [50]: scipy.misc.imsave('sample_resize.jpg', img_resized)
In [51]: ls sam
sample.jpg          sample_resize.jpg
In [51]: ls sample
```

# Little Demo


input text: black hair blue eyes



input text: pink hair green eyes



input text: green hair green eyes



input text: blue hair red eyes



# Tips for training

- Discriminator output:
  - (real img, right text): 1
  - (fake img, right text): 0
  - (real img, wrong text): 0
  - (wrong img, right text): 0
- Different objective function
  - Wasserstein GAN (WGAN)
  - Least Squares GAN (LSGAN)
  - Boundary-Seeking GAN (BGAN)

# Wasserstein GAN

- Earth-Mover (Wasserstein-1) distance

$$\begin{aligned} W(P_r, P_g) &= \inf_{\gamma \in \Pi(P_r, P_g)} E_{(x,y) \sim \gamma} [\|x - y\|] \\ &= \sup_{\|f\|_L \leq 1} \{ E_{x \sim P_r} [f(x)] - E_{x \sim P_\theta} [f(x)] \} \end{aligned}$$

EM distance is Continuous  
JS-divergence is not

- The optimization problem is thus :

$$\begin{aligned} \max_{w \in W} \{ E_{x \sim P_r} [f_w(x)] - E_{x \sim P_\theta} [f_w(x)] \} \\ = \max_{w \in W} \{ E_{x \sim P_r} [f_w(x)] - E_{z \sim P_z(z)} [f_w(g_\theta(z))] \} \end{aligned}$$

# Wasserstein GAN

- For Discriminator

The output of D is thus not probability anymore !  
The D loss turn to be a measure of distance

$$\max_D \{ E_{x \sim P_r} [D(x)] - E_{z \sim P_z(z)} [D(G(z))] \}$$

- For Generator

$$\max_G \{ E_{z \sim P_z(z)} [D(G(z))] \}$$

# Wasserstein GAN

Implementation Notes :

- Do not apply sigmoid at the output of D
- Clip the weight of D
- Use RMSProp instead of Adam
- Train more iteration of D (the paper use 5)

# Least Squares GAN

- For Discriminator


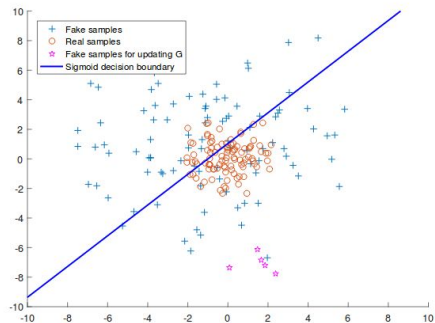
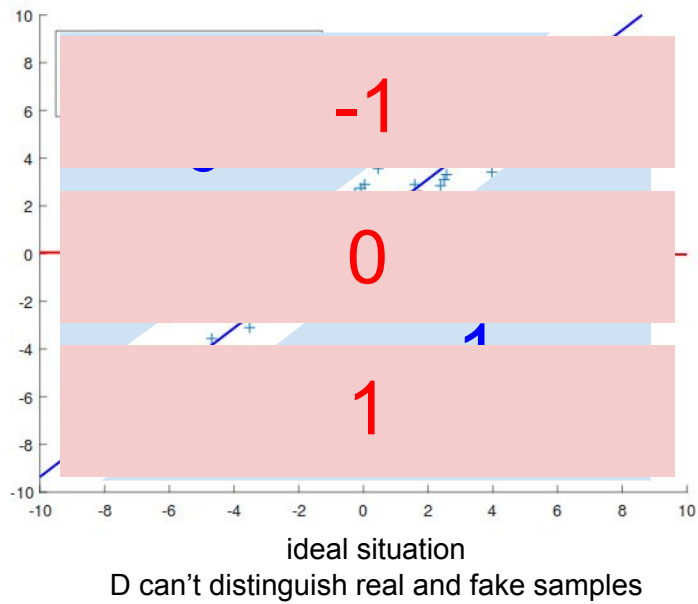
$$\min_G \left\{ \frac{1}{2} E_{x \sim p_{data}(x)} [(D(x) - \overset{1}{b})^2] + \frac{1}{2} E_{z \sim p_z(z)} [(D(G(z)) - \overset{-1}{a})^2] \right\}$$

- For Generator



$$\min_G \frac{1}{2} E_{z \sim p_z(z)} [(D(G(z)) - \overset{0}{c})^2]$$



# Least Squares GAN



Drag to decision boundary



# Least Squares GAN

- For Discriminator

$$\min_G \left\{ \frac{1}{2} E_{x \sim p_{data}(x)} [(D(x) - b)^2] + \frac{1}{2} E_{z \sim p_z(z)} [(D(G(z)) - a)^2] \right\}$$

- For Generator

$$\min_G \frac{1}{2} E_{z \sim p_z(z)} [(D(G(z)) - c)^2]$$

# Least Squares GAN

Implementation Notes :

- Do not apply sigmoid at the output of D

# Boundary-Seeking GAN

Optimum Discriminator :

$$D^*(x) = \frac{P_{data}(x)}{P_{data}(x) + P_g(x)}$$



$$P_{data}(x) = P_g(x) \frac{D^*(x)}{1 - D^*(x)}$$



For imperfect discriminator :

$$\tilde{P}(x) = \frac{1}{Z} P_g(x) \frac{D(x)}{1 - D(x)}$$

**Z : normalize term**

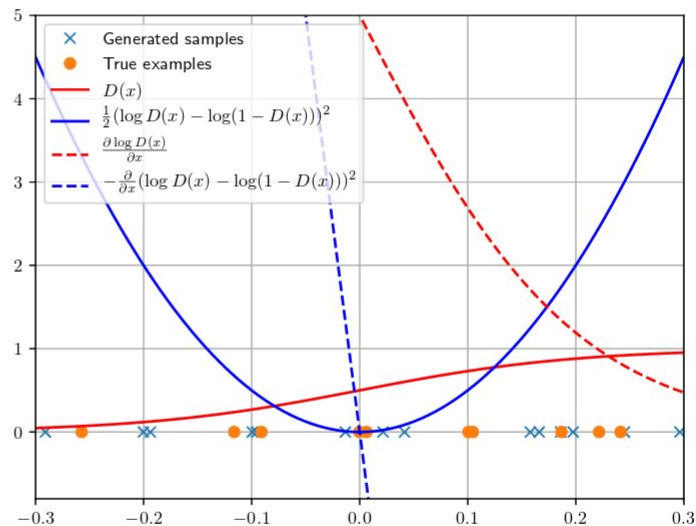

# Boundary-Seeking GAN

- For Generator

$$\min_G E_{z \sim P_z(z)} [(\log(D(G(z))) - \log(1 - D(G(z))))^2]$$

- Only change the loss function of G

# Boundary-Seeking GAN




# Submission and Grading

# Homework 3 package

- Anime Dataset:

- training data: 33.4k (image, tags) pair
- faces/, tags.csv, sample\_testing.txt



blue eyes  
red hair  
short hair

- training tags file format

- img\_id <comma> tag1 <colon> #\_post <tab> tag2 <colon> ...

```
1 0,touhou:17705 |chen:423 |moneti daifuku :60 |animal ears:12241 |catgirl:4903 |
2 1,touhou:17697 |onozuka komachi:224 |shikieiki yamaxonadu:217 |$
3 2,original:25774 |blonde hair:25457 |doll:1040 |dress:16585 |pink eyes:3896 |ta
4 3,amagi brilliant park:111 |musaigen no phantom world:39 |nichijou:142 |kawakam
```

tags.csv

- testing text file format

- testing\_text\_id <comma> testing\_text

```
1 1,blue hair blue eyes
2 2,blue hair green eyes
3 3,blue hair red eyes
4 4,green hair blue eyes
```

sample  
testing\_text.txt

- testing text only includes **'color hair'** and **'color eyes'**, only alphabetic char involved.

- Data download link:

- <https://drive.google.com/open?id=0BwJmB7aIR-AvMHEtczZZN0EtdzQ>
- If you want to do something cool beyond generating faces, mail us. We will give you original images



# Submission

- Only **Python** with **Tensorflow r1.0** for GAN
  - You are allowed to use any tool to process text
  - You are allowed to use any data you collect
- Deadline: **5/25(Thu.) 23:59:59 (UTC+8)**
- **MLDS2017/hw3** should contain following files
  - **run.sh train.py, (pre-)trained\_model, generate.py, samples/, report.pdf**
  - If some files are too big, upload to your cloud and download them when running your run.sh
- TAs will run your run.sh to generate images given a text
  - `bash run.sh [testing_text.txt]`
  - run.sh must output in **10 minutes**.

# Submission on Github

- Only one branch **master** is needed
- **master** stores the model by **using GAN structure**
- Remember to put your **pre-trained models or download scripts** so that we can run your code successfully

# Output Format Requirement

- The generated images should be in Directory **samples/**
  - make sure it's **empty** before we run your code
- Each generated image must be **64 x 64** in size
  - please resize all training images to 64 x 64 before training the model
- For each input text, you must generate **5 images**
- Generated img should be named as  
**"sample\_(testing\_text\_id)\_(sample\_id).jpg"**
- Example:

```
andy@andy-All-Series|x86_64:samples:4$ ls
sample_1_1.jpg sample_1_3.jpg sample_1_5.jpg sample_2_2.jpg sample_2_4.jpg
sample_1_2.jpg sample_1_4.jpg sample_2_1.jpg sample_2_3.jpg sample_2_5.jpg
```

# 組別互評

- We will put your generated images in the grading platform
- Link will be sent to your mail after HW deadline
- Answer **2** scores for each image
  - How the image fits the text
  - How the image looks real
- Scores should be integer from 1 to 5
  - 1 to 5 corresponding to (super bad, bad, average, good, super good)
- You may score your results, so be fair when your are scoring :)

# 組別互評

- Separate scores with a comma (score for matching text, score for reality)
- Example:

3 → Gray hair green eyes



4, 5

Ok ✓

press ENTER

# What report should cover?

- Environment (1%)
  - Ex. OS, CPU, GPU, Memory, libraries you used and version, etc.
- Model description(3%)
  - Must include model structure, objective function for G and D
- How do you improve your performance (5%)
- Experiment settings and observation (5%)
- Team division (1%)
- No more than 5 pages
- Please written in Chinese (unless you don't know how to type Chinese)

# Grading Policy (20%)

- Wrong output format will not be graded
- Report (10%)
- Code (5%)
  - You will be scored only if you use GAN and output results in 10 minutes
- ~~Baseline~~
- 組別互評 (5%)
- Bonus
  - 組別互評 top 3 (5%, 3%, 2%)
  - Generate images beyond faces and show results in report (10%)

# Other Policy

- Late policy: 30% off per day late afterwards.  
[Delay form will be announced afterwards]
- No plagiarism is allowed.